

# ELIXIRCONF™ 2017



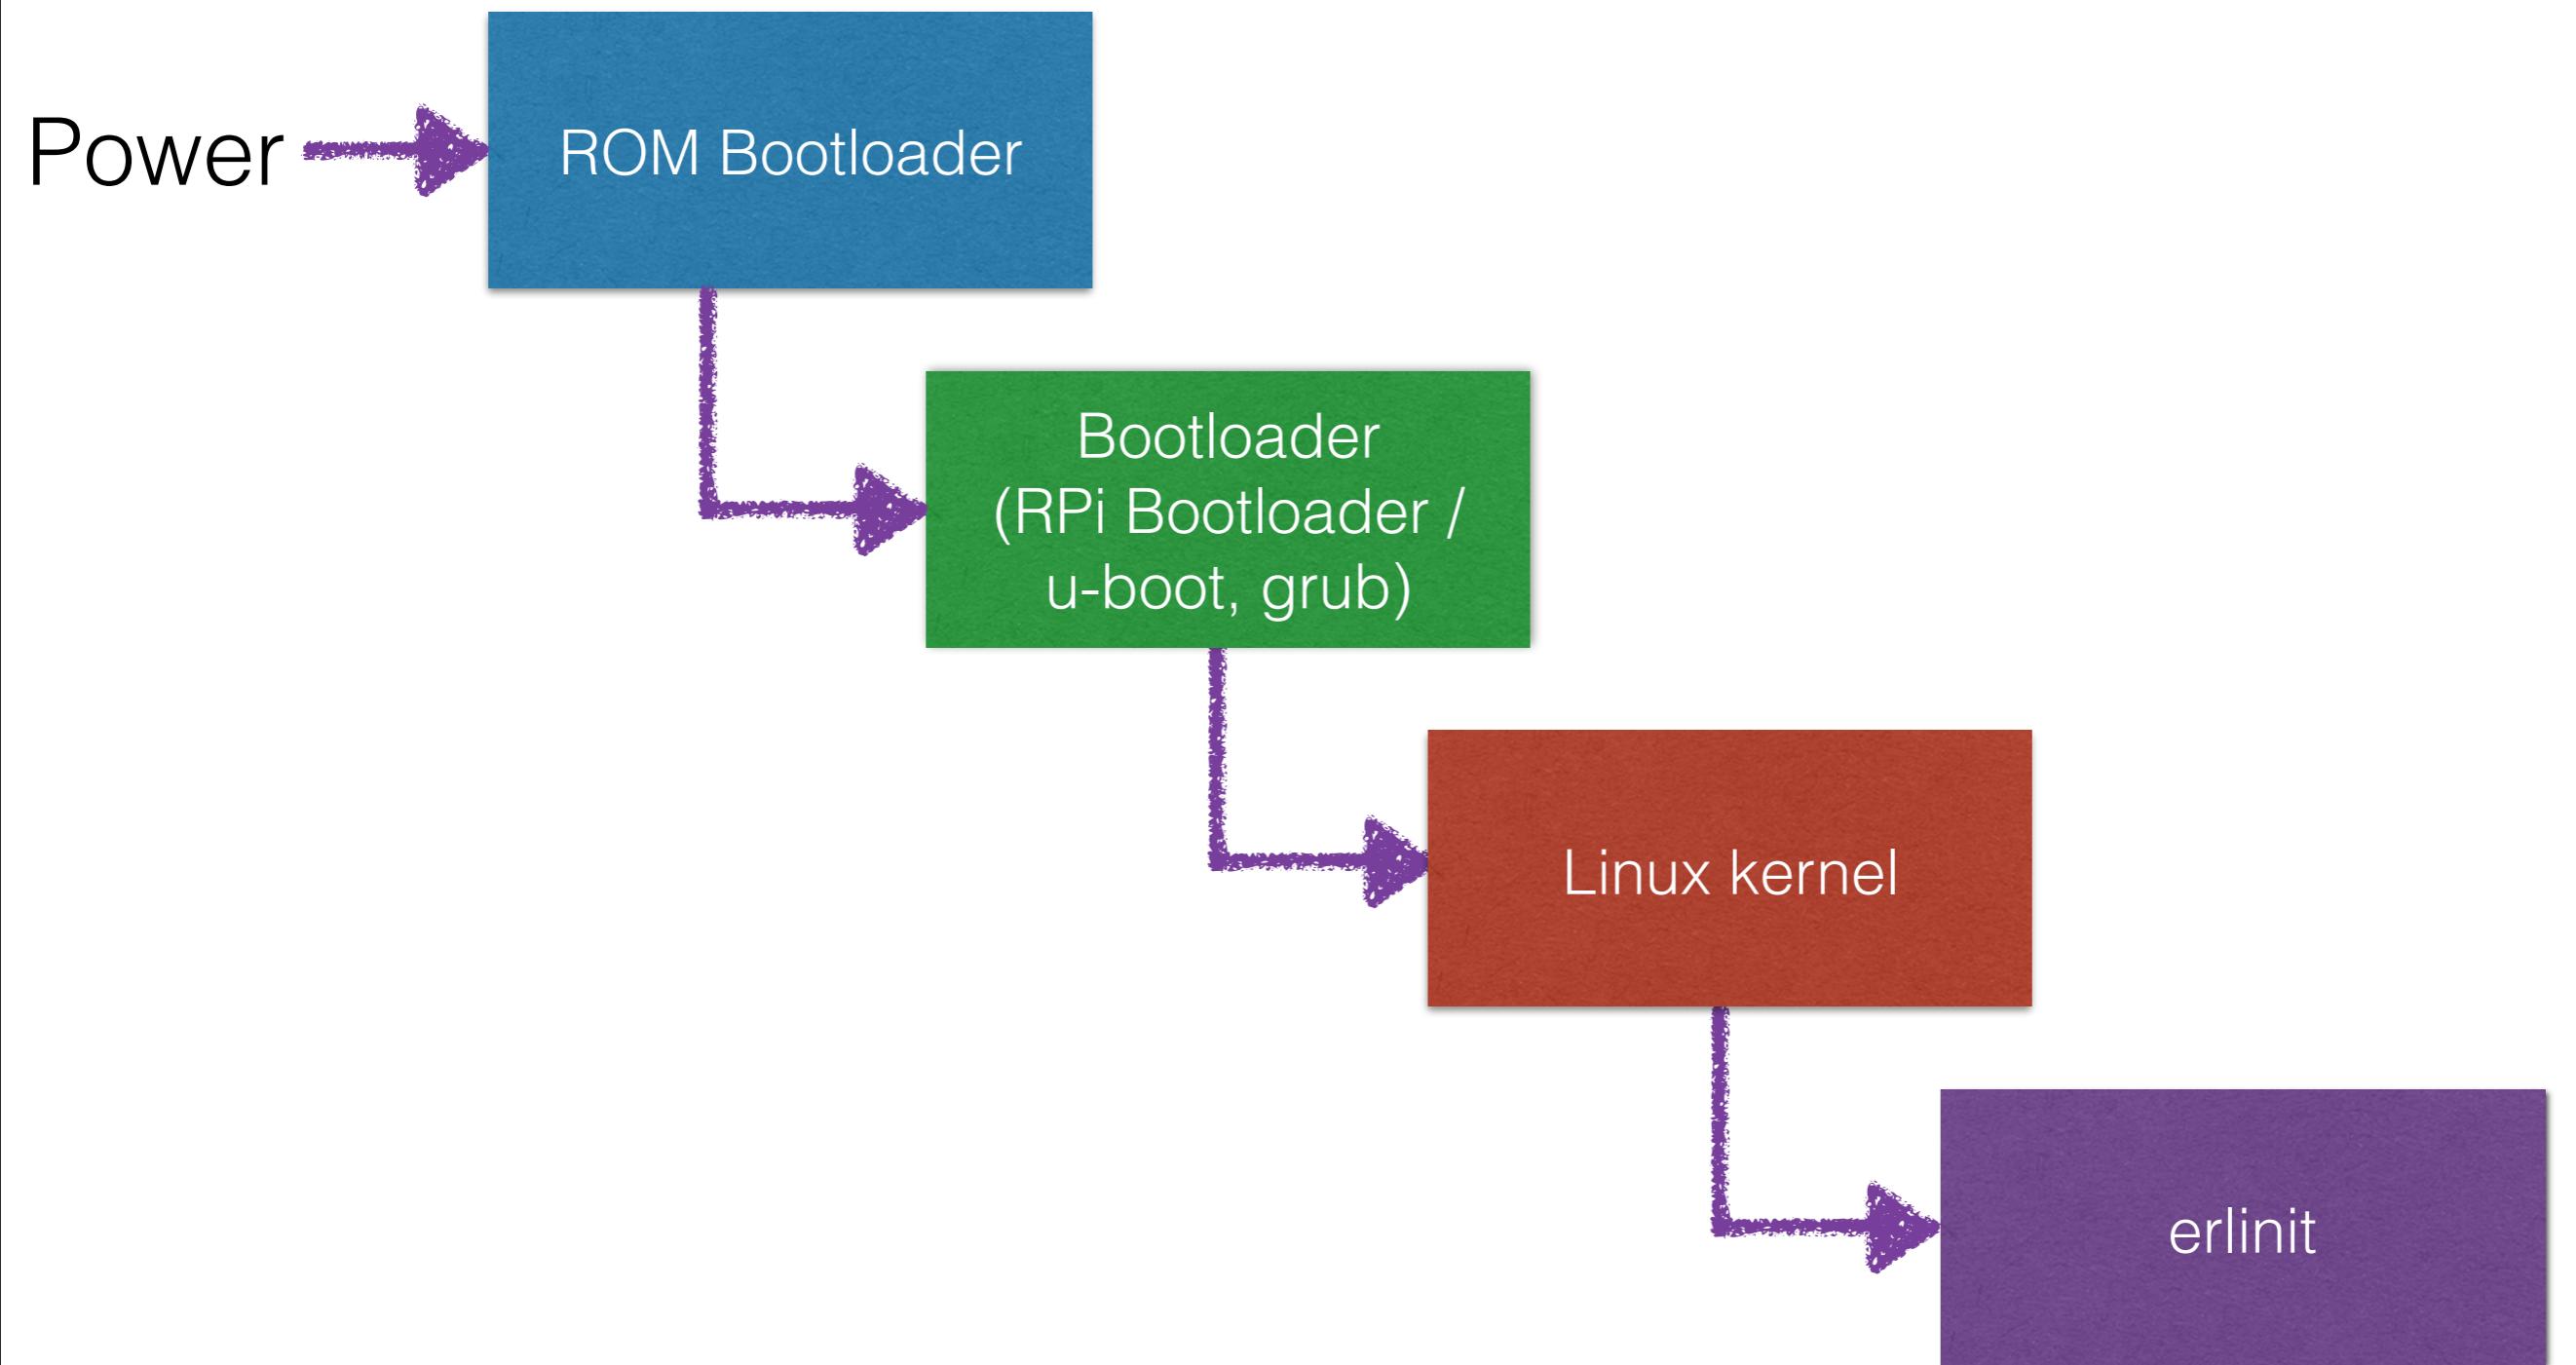
## Nerves Training Part 7: Advanced Topics

# Topics

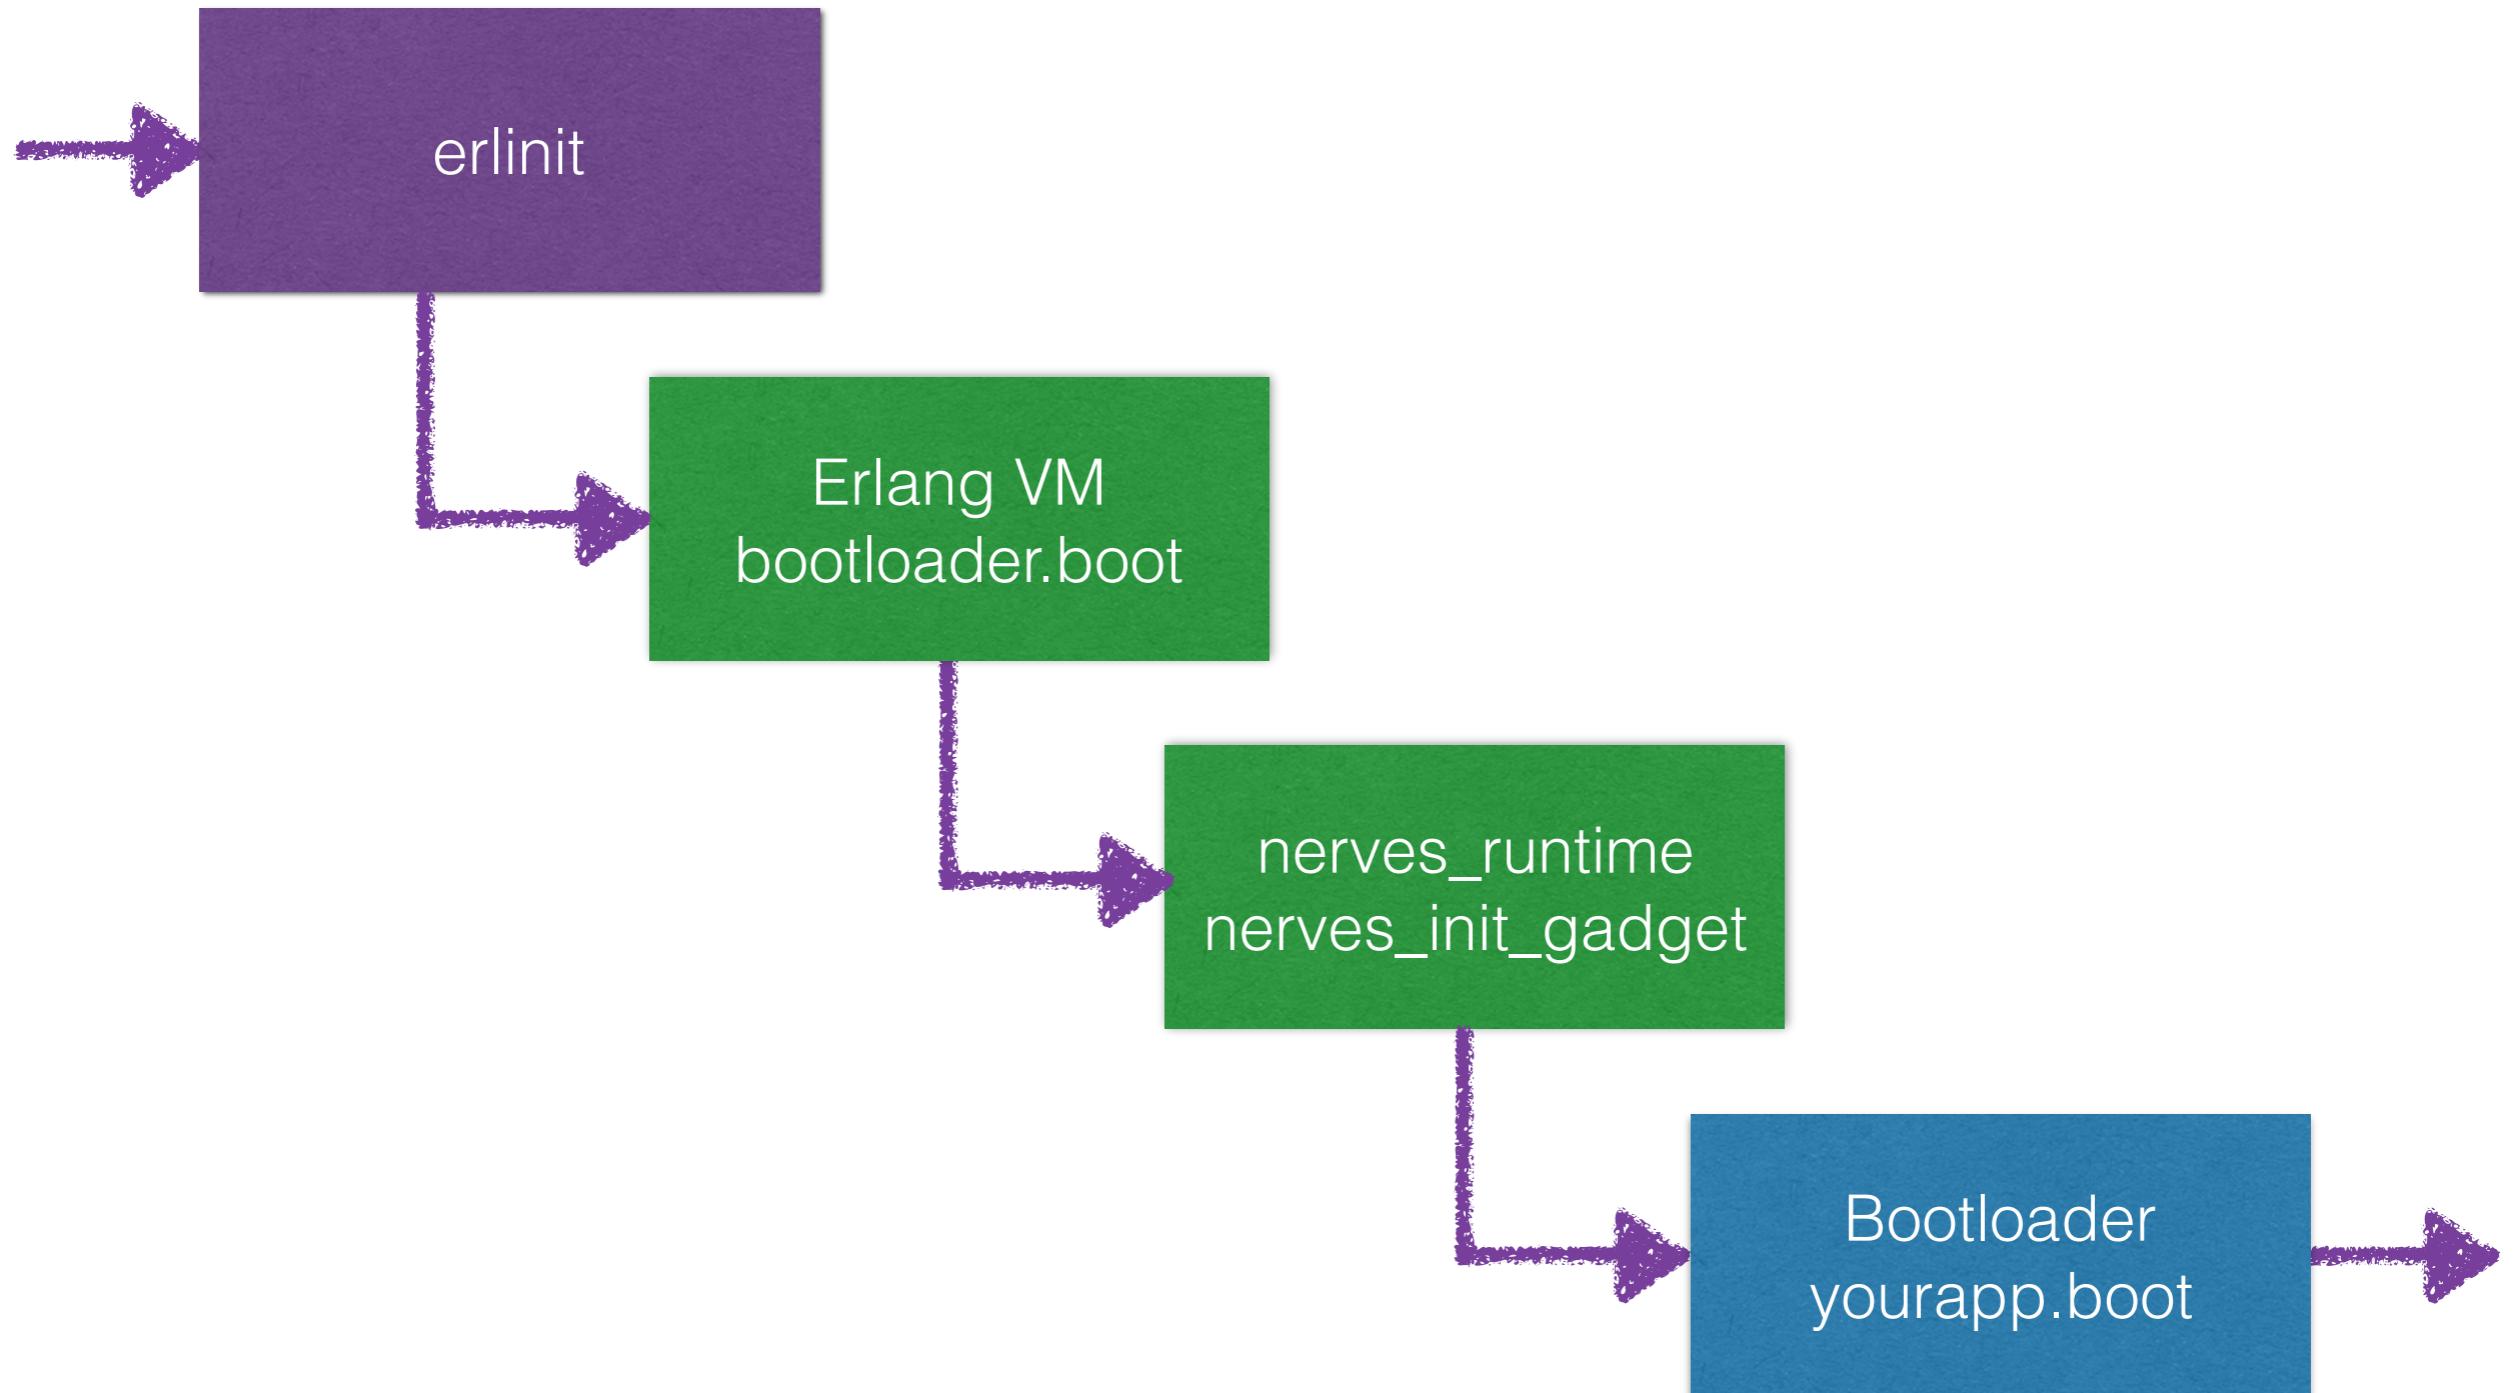
- Device initialization
- Storage layout and firmware updates
- Nerves and hard real-time
- Interfacing with C
- Security
- Going forward

# Initialization

# Initialization



# Initialization Part 2



# nerves\_runtime

- Handles common initialization for all Nerves devices
- Initializes the application data partition if necessary
  - Formats as ext4
  - Mounts as /root by default
- Starts monitoring device insertion and removal

# erlinit

- Replacement for /sbin/init that starts the Erlang virtual machine
- Basic initialization of the Linux user land
  - Loopback network connection
  - Mounts /tmp, /proc, /sys
  - Configures the tty
- Configuration stored in /etc/erlinit.conf
  - Can be overridden by passing parameters to the Linux kernel from the bootloader

# erlinit debugging options

- `--verbose`
- `--hang-on-exit`
  - Useful to capture error messages when the VM exits
- `--run-on-exit /bin/sh`
  - Drop into a shell if the VM exits
  - Exiting the shell reboots or hangs
- `--warn-unused-tty`
  - erlinit will tell you what options to pass it to use the shell on the terminal you're looking at

# erlinit features

- Mount filesystems
- Configure a unique hostname
  - --hostname-pattern and --uniqueid-exec
  - Nerves uses fhunleth/boardid to read serial numbers
- Wrap the launch of the Erlang VM in another program
  - --alternate-exec
  - Perform some custom system-specific initialization that can't be done in Erlang or Elixir
  - Capture the terminal with dtach to route it to a GUI
- Run the Erlang VM as a regular user

# erlinit pitfalls

- Running shell scripts to initialize the system
  - Move initialization to Elixir to take advantage of OTP supervision
  - Be aware that Linux daemon processes don't supervise easily
- Assuming writable filesystems always can be mounted
  - Failures happen – must be handled in Elixir
  - erlinit can't report errors so Elixir must check

# Topics

- Device initialization
- Storage layout and firmware updates
- Nerves and hard real-time
- Security
- More resources

# Storage

# Storage Basics



Raw NAND Flash



MicroSD



eMMC

- Almost all embedded devices use NAND Flash now since it's large and cheap
- MicroSD cards and eMMC contain microcontrollers that handle wear leveling, bad block handling, etc.

# Basic Flash Layout

- Master Boot Record (MBR)
  - Table of contents
- Boot partition readable by ROM bootloader
  - Linux, device tree and other files
- Root filesystem
  - Almost everything under "/"
  - Normally writable

Master Boot Record

Boot partition (FAT)

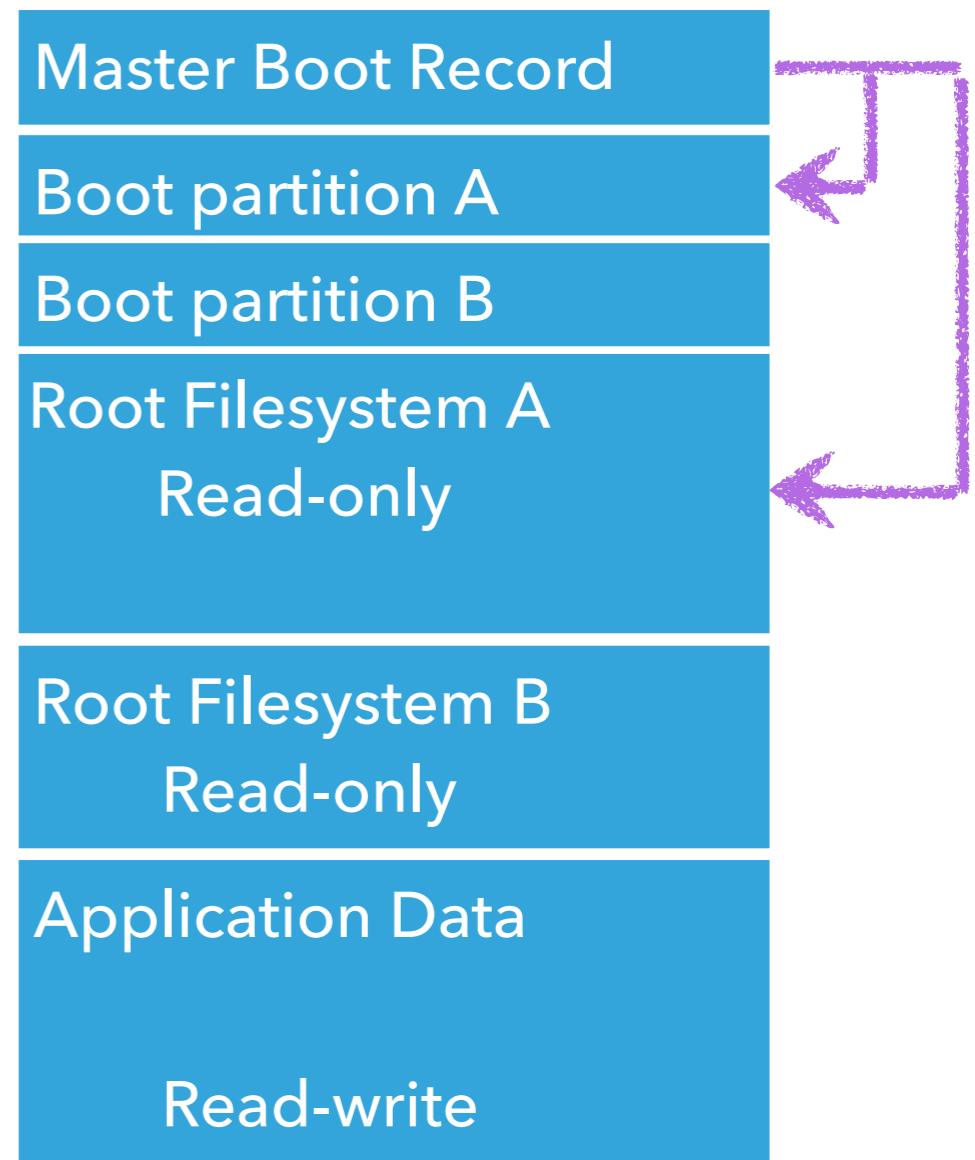
Root filesystem

# Common Filesystems

- FAT8, FAT16 and FAT32 - the filesystems that you know and love from DOS and Windows
  - Very simple and widely supported even though old tech
  - Get corrupted very easily
- ext4 - de facto Linux filesystem
  - Relatively fast and supports journals for failure recovery
  - Not so well supported outside of Linux world
- SquashFS - compressed read-only filesystem
  - Common in embedded systems and Live CDs
  - Often combined with UnionFS
- UBIFS - If you use raw NAND Flash

# Flash Layout with Nerves

- Pair boot and root filesystems
  - Updates go to the unused one
  - Swap active when done
- Read-only root filesystem
- If data partition is corrupt reformatting it returns system to a usable state



# Root filesystem overlay

- Needed to update files outside of Elixir
  - Linux program configuration files in /etc
  - erlinit.config
  - Workaround read-only filesystem with symlinks to the application partition
- See *config/config.exs* for example

# Failure scenarios

- Root filesystem corruption due to ill-timed disk op
  - Difficult thanks to read-only filesystem
- Ill-timed power failure or crash on firmware update
  - Small window: A/B swap is final operation and nearly atomic
- Bad firmware applied
  - Cryptographic hashes to detect corruption
  - Verification of firmware metadata before applying it

# Handling Faulty Firmware

- Development scenarios
  - Use Elixir bootloader to initialize everything needed for firmware updates
  - Take advantage of `nerves_network` configuration cleanup when processes die

# Handling Faulty Firmware 2

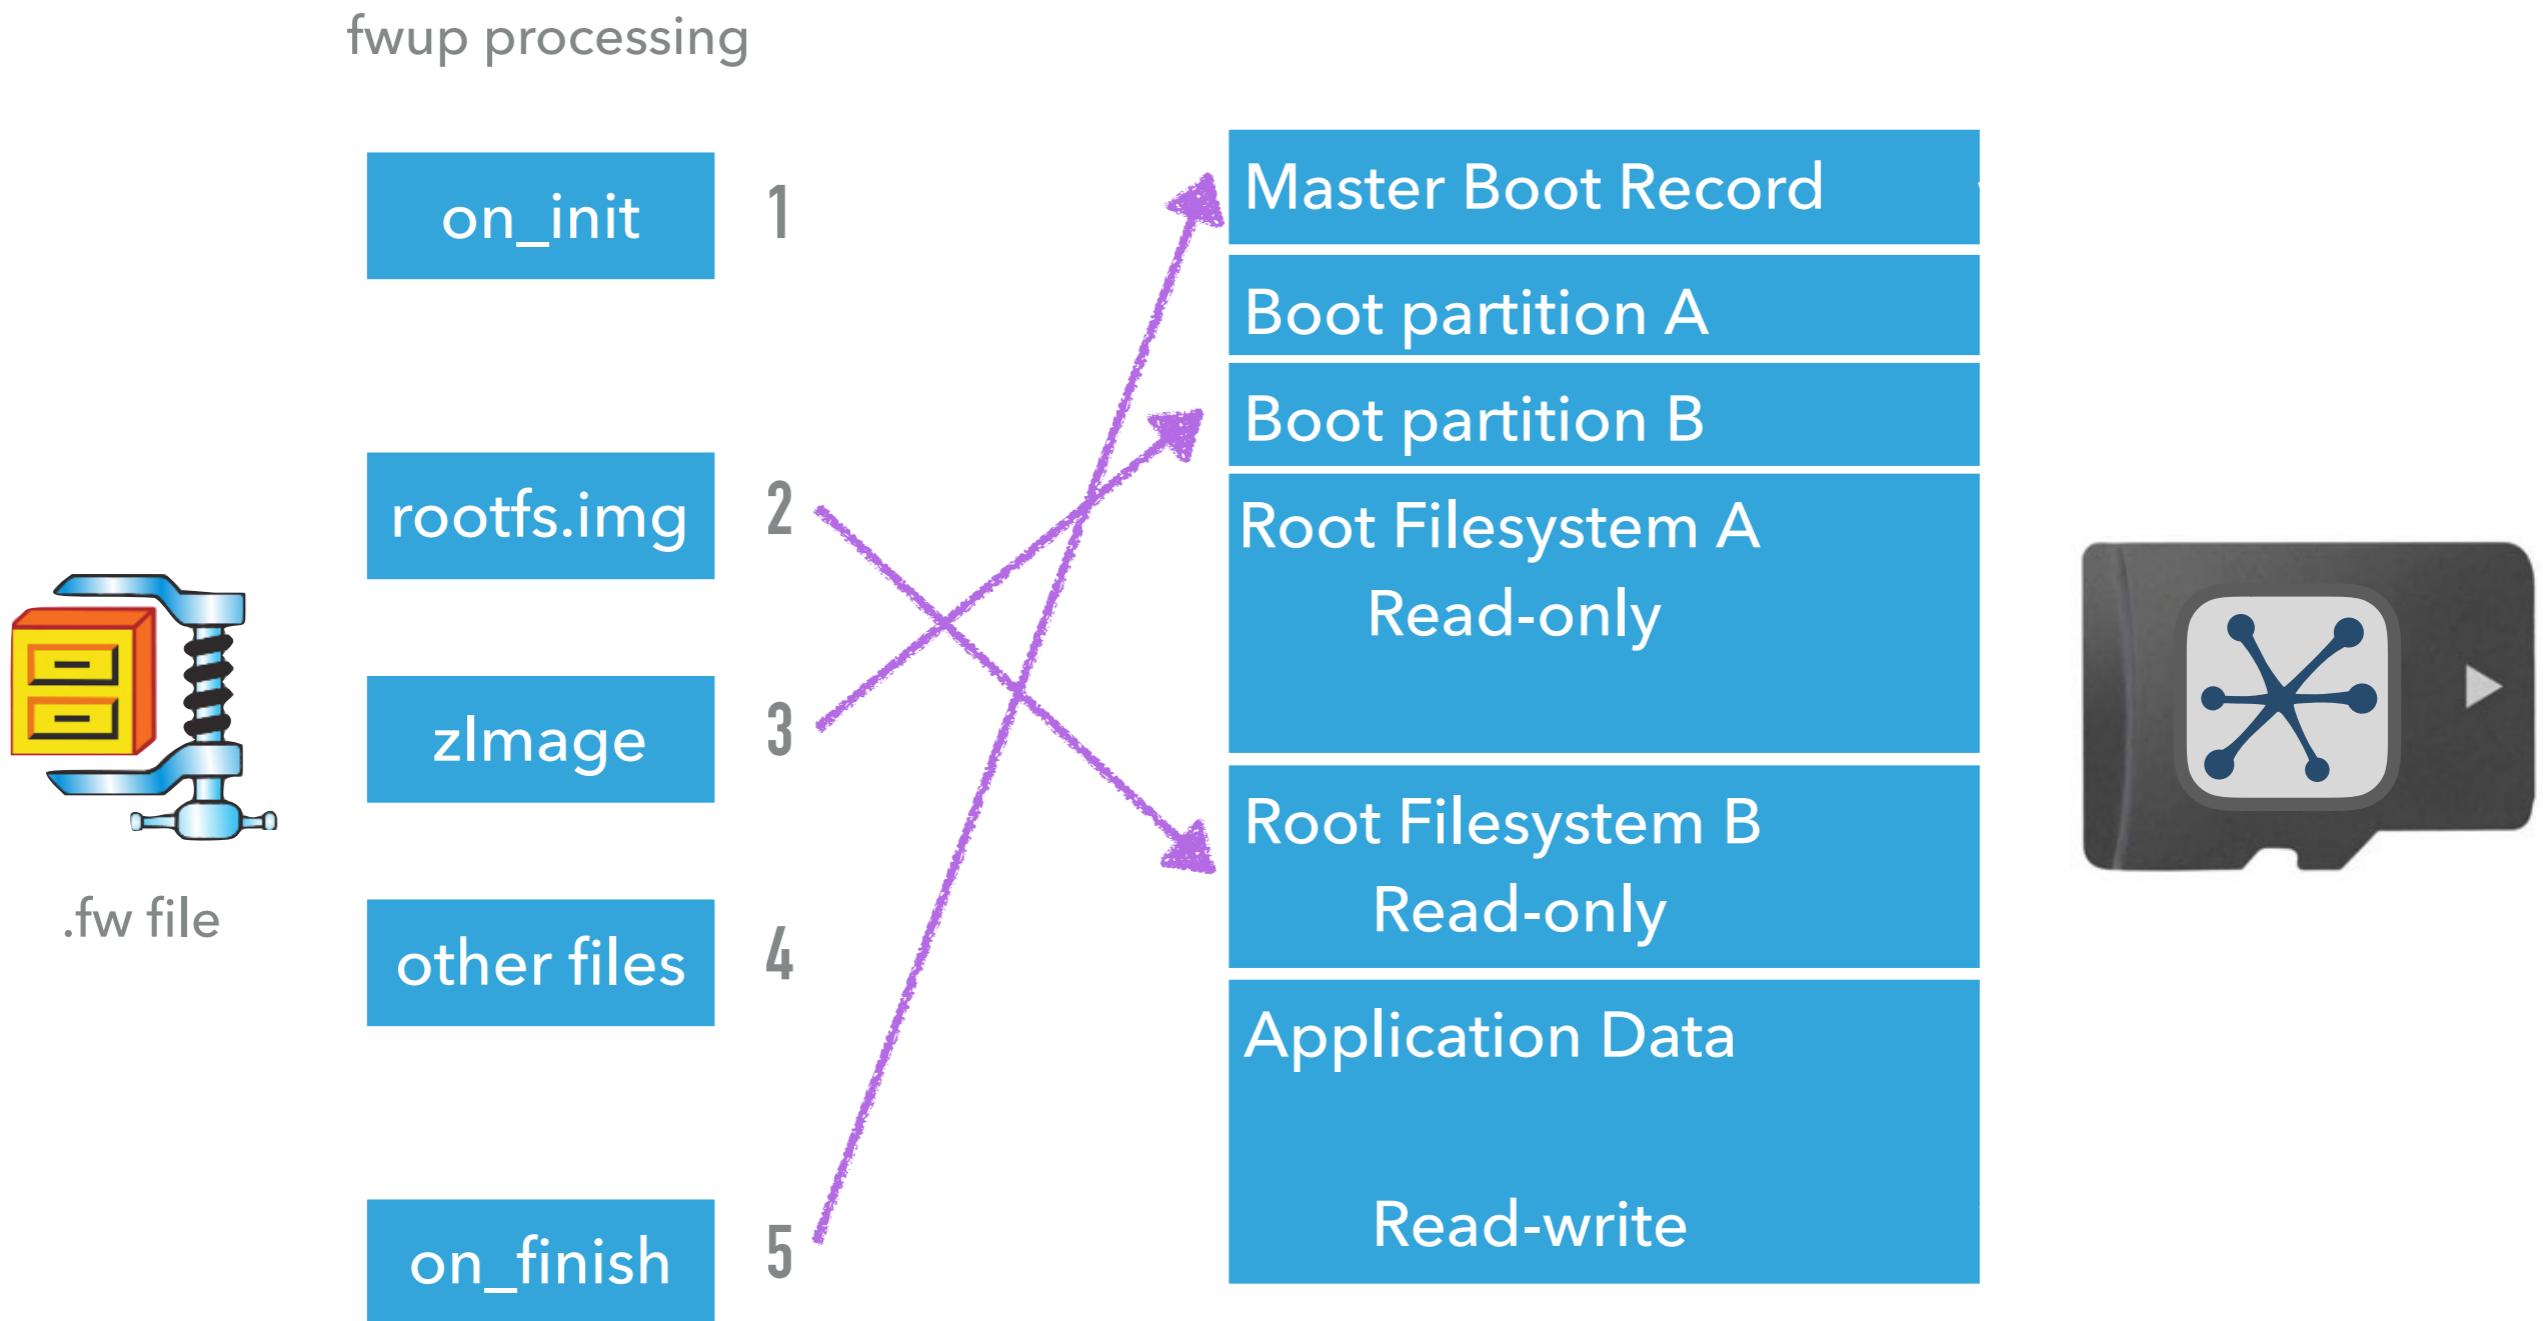
- Verify and fallback
  - Allow one boot of new firmware
  - If "valid" bit isn't set, then next reboot reverts to old firmware
- Safe mode
  - Dedicated partition that runs a firmware update application
  - Detect button press (or other input) at start to enter this mode

# fwup

- Firmware update packaging and application
- Packages
  - Zip-formatted archives
  - Metadata
  - All data files protected by cryptographic hashes
- Packages can be cryptographically signed
- Very simple scripts supported (lack of functionality is a feature)



# fwup processing



# Reasons to Modify fwup.conf

- Raspberry Pi bootloader configuration changes
  - New device tree overlays for HATs
  - Change UART config, GPU parameters, etc.
  - Change Linux kernel parameters
- Implement fallback and safe boot
- Modify root filesystem and data partition sizes

# Nerves Firmware and System Metadata

# Nerves.Runtime.KV

- Simple key-value store for system and firmware metadata
- Not intended to be written frequently
- Requires Nerves system support
  - U-boot formatted kv store not in any partition
  - fwup.conf can modify keys
  - Requires uboot-tools (*fw\_printenv* and *fw\_setenv*)

# Nerves.Runtime.KV Uses

- Information about running firmware
  - Useful for bug reports and new release detection
  - Verify firmware update compatibility (versions, architecture, product, etc.)
- Device provisioning information
  - Calibration
  - Manufacturing build configuration

# Try it out

```
iex> Nerves.Runtime.KV.get_all_active
%{"nerves_fw_application_part0_devpath" => "/dev/mmcblk0p3",
 "nerves_fw_application_part0_fstype" => "ext4",
 "nerves_fw_application_part0_target" => "/root",
 "nerves_fw_architecture" => "arm",
 "nerves_fw_author" => "The Nerves Team",
 "nerves_fw_description" => "",
 "nerves_fw_misc" => "",
 "nerves_fw_platform" => "rpi0",
 "nerves_fw_product" => "starter",
 "nerves_fw_vcs_identifier" => "",
 "nerves_fw_version" => "0.1.0"}
```

Works on any of the training examples

# Nerves and Hard Real-time

# Real-time

- How critical is response time to the correctness of a program?
- Real-time systems provide guarantees
  - Hard real-time - Deadlines can never be missed
  - Soft real-time - Missed deadlines are undesirable, but not fatal
  - Not real-time - No timeliness guarantees are made

# Cooperative Scheduling

- Any system using an event loop
  - Node.js
  - Python Twisted
  - C++ Qt framework
- No guarantees that a long running computation doesn't stall the event loop
- Does that 3rd party package I just pulled in behave nicely?
- Not a real-time system

# Scheduling in Elixir

- Preemptive multi-tasking
- Each process gets a fixed number of “reductions”
  - Function calls
  - Garbage collects
  - Sending messages
- Processes are scheduled round robin
- Four queues (:low, :normal, :high, :max) but priorities other than :normal are rarely used
- Soft real-time

# Real-time Operating Systems

- Preemptive multi-tasking
- Often priority based schedulers, but other options available
- Can handle priority inversion (when a high priority process needs something from a low priority one)
- Run in limited resource environments
- Fewer features than general purpose operating systems
- Support hard real-time applications
- See RTEMS usage with GRiSP

# Hard Real-time Tasks

- Collecting data from streaming sensors
- Sensor/actuator feedback loops
- Low level communication with hardware devices (e.g., LED strips)
- Example: Accelerometer
  - Acceleration data reported at 100+ Hz
  - If sample missed, it's lost - no retransmissions
- Example: PID motor controller
  - Feedback from RPM sensor used to control current to a motor
  - Control algorithm tuning depends on algorithm w/ precise timing

# Mix of tasks

Hard real-time ← → Soft real-time

Touchscreen UI

Data collection

Web interface

Control logic

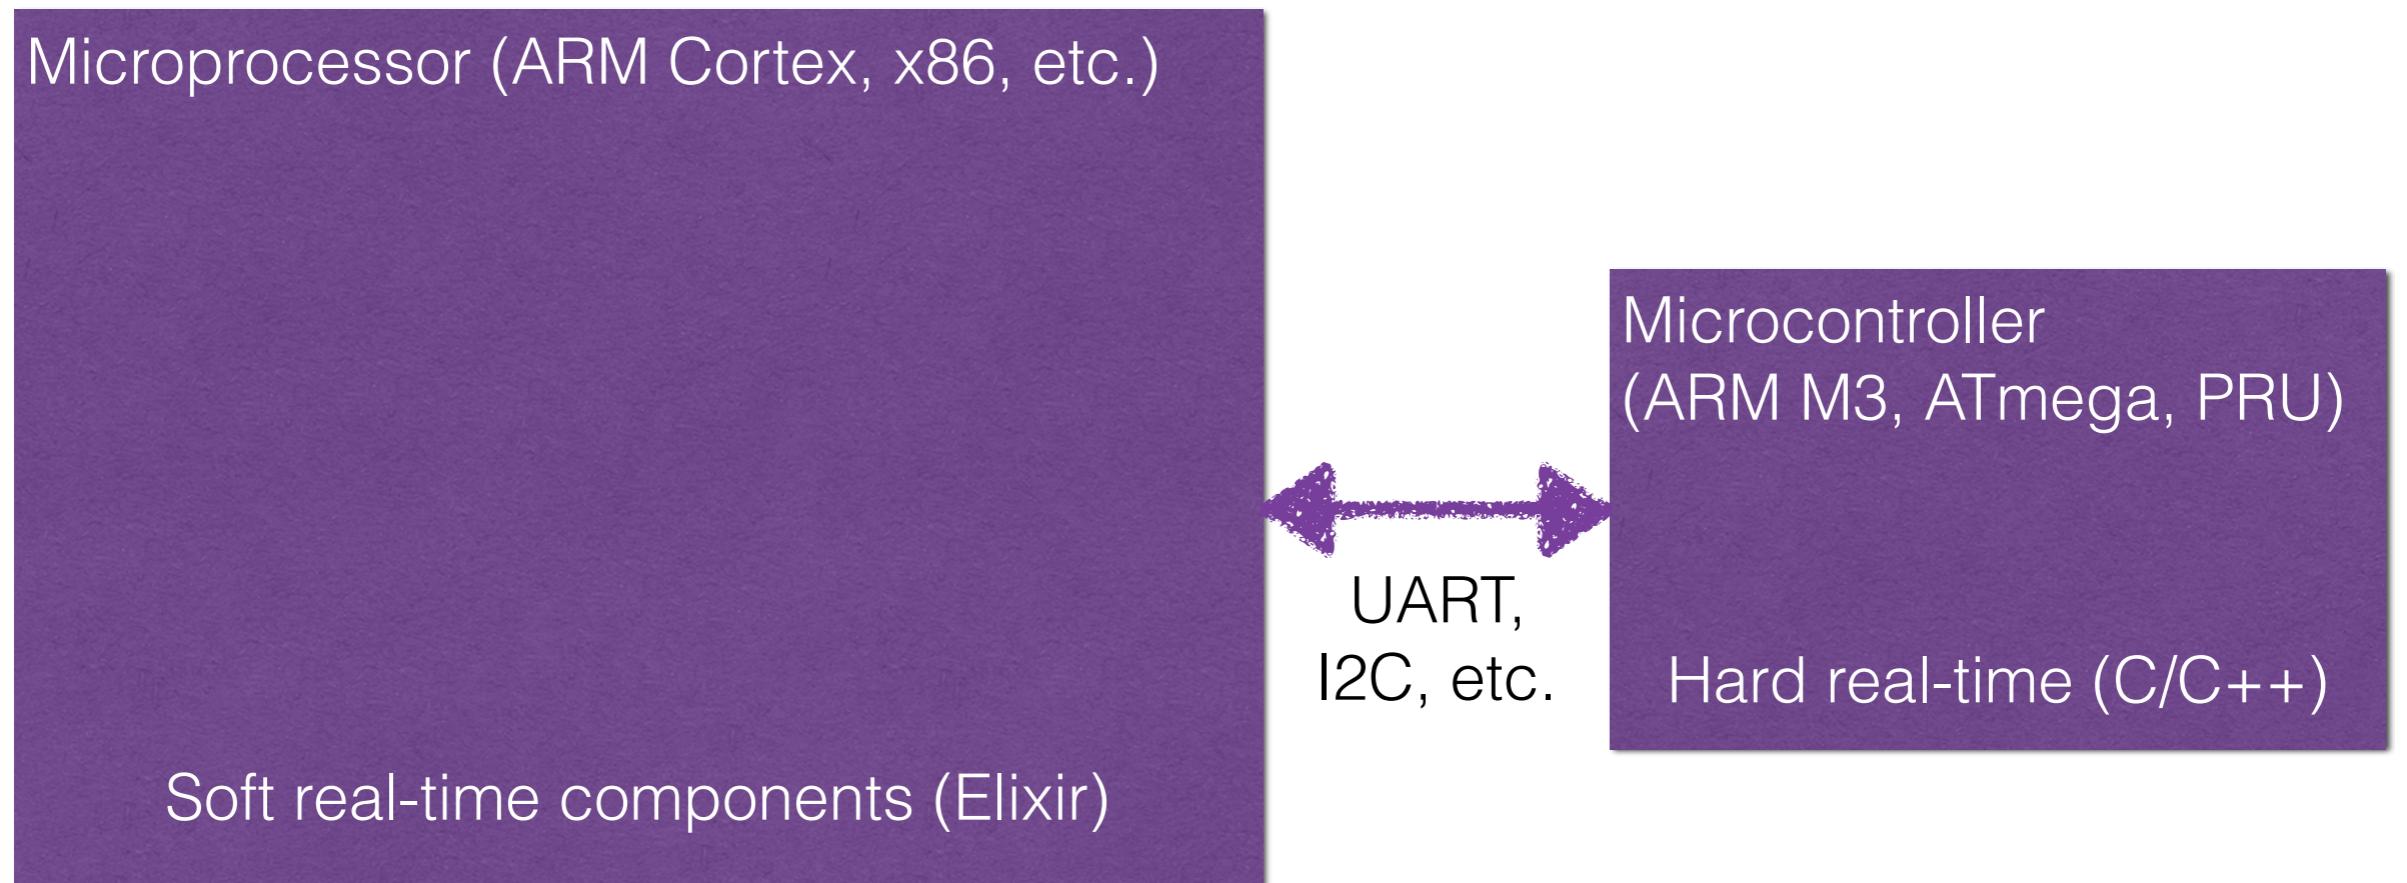
Software update

Actuator control

System configuration

Diagnostics

# Common Architecture



Segregating hard real-time tasks can simplify both verification of real-time constraints and overall system complexity.

# Interfacing with C

# Elixir is not a Systems Programming Language

- Can't read or write to special files like device files
- No ioctl()
- No mmap()
- Low level OS parts of C standard library generally missing (e.g., signals)

Not really a problem, since providing this many of these features inside the Erlang VM might compromise soft real-time and stability

# Interfacing to C (ports)

- Start a program and communicate via stdin/stdout
- OS process isolation from the Erlang VM
- Debug standalone or via calls from Elixir
- `erl_interface` provides some help with marshaling and unmarshaling Erlang terms in C (optional)

# Example Port Protocol

message\_size

{request, args}



To port process

```
msg = {request, args}  
send myport, {self, {:command, :erlang.term_to_binary(msg)}}
```

message\_size

return\_value

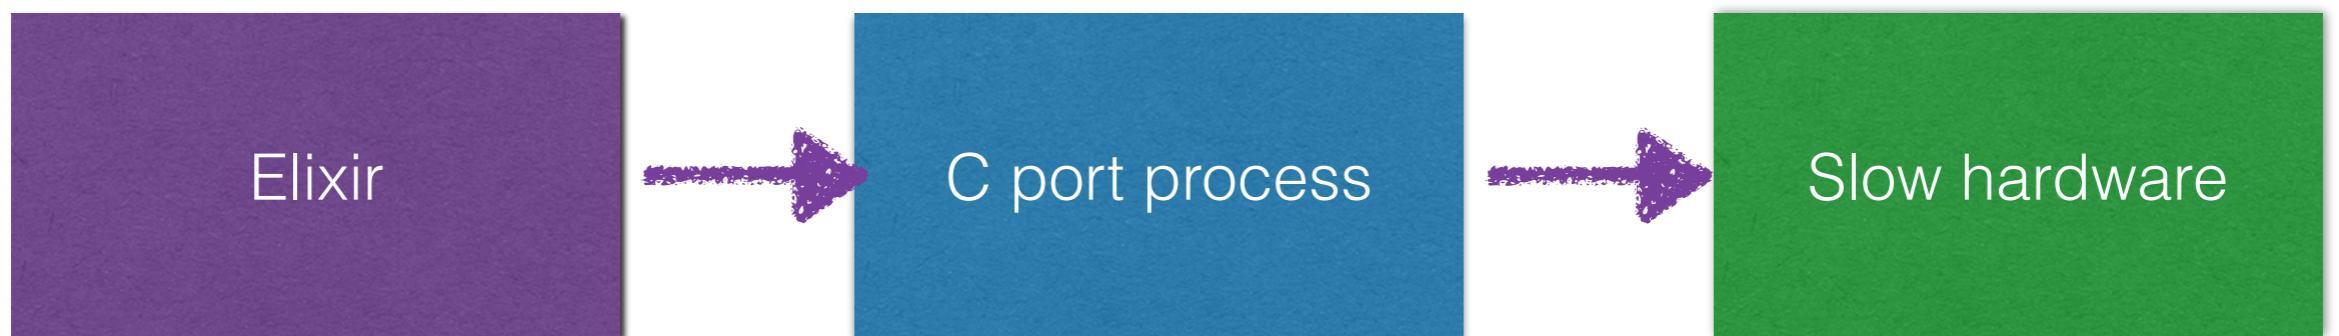


From port process

```
receive do  
  {_, {:data, return_value}} ->  
    result = :erlang.binary_to_term(return_value)  
    Logger.debug "Got #{inspect result}"  
end
```

# Queuing and Back Pressure

- Without back pressure, requests queue between Elixir code and the port process
- Easy solution: wait for a response for each request



# Limitations

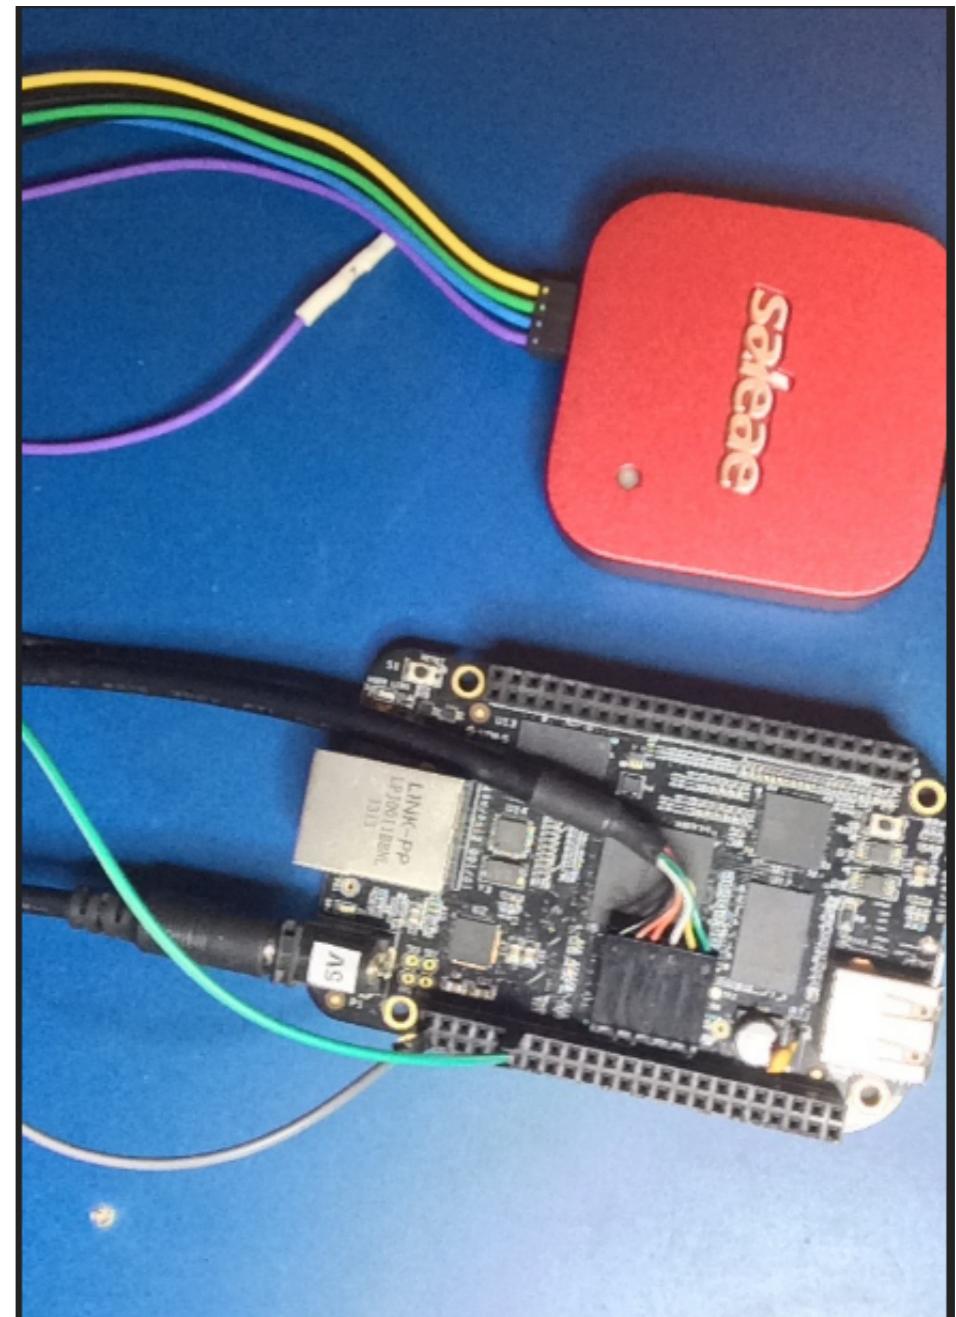
- Erlang ports lack common operations on child processes
  - Can't send EOF - affects nearly all Unix data processing programs
  - Can't send signals
- Processes scheduled by the OS scheduler (not Erlang's)
- Misbehaving C processes can easily stick around even when Erlang port

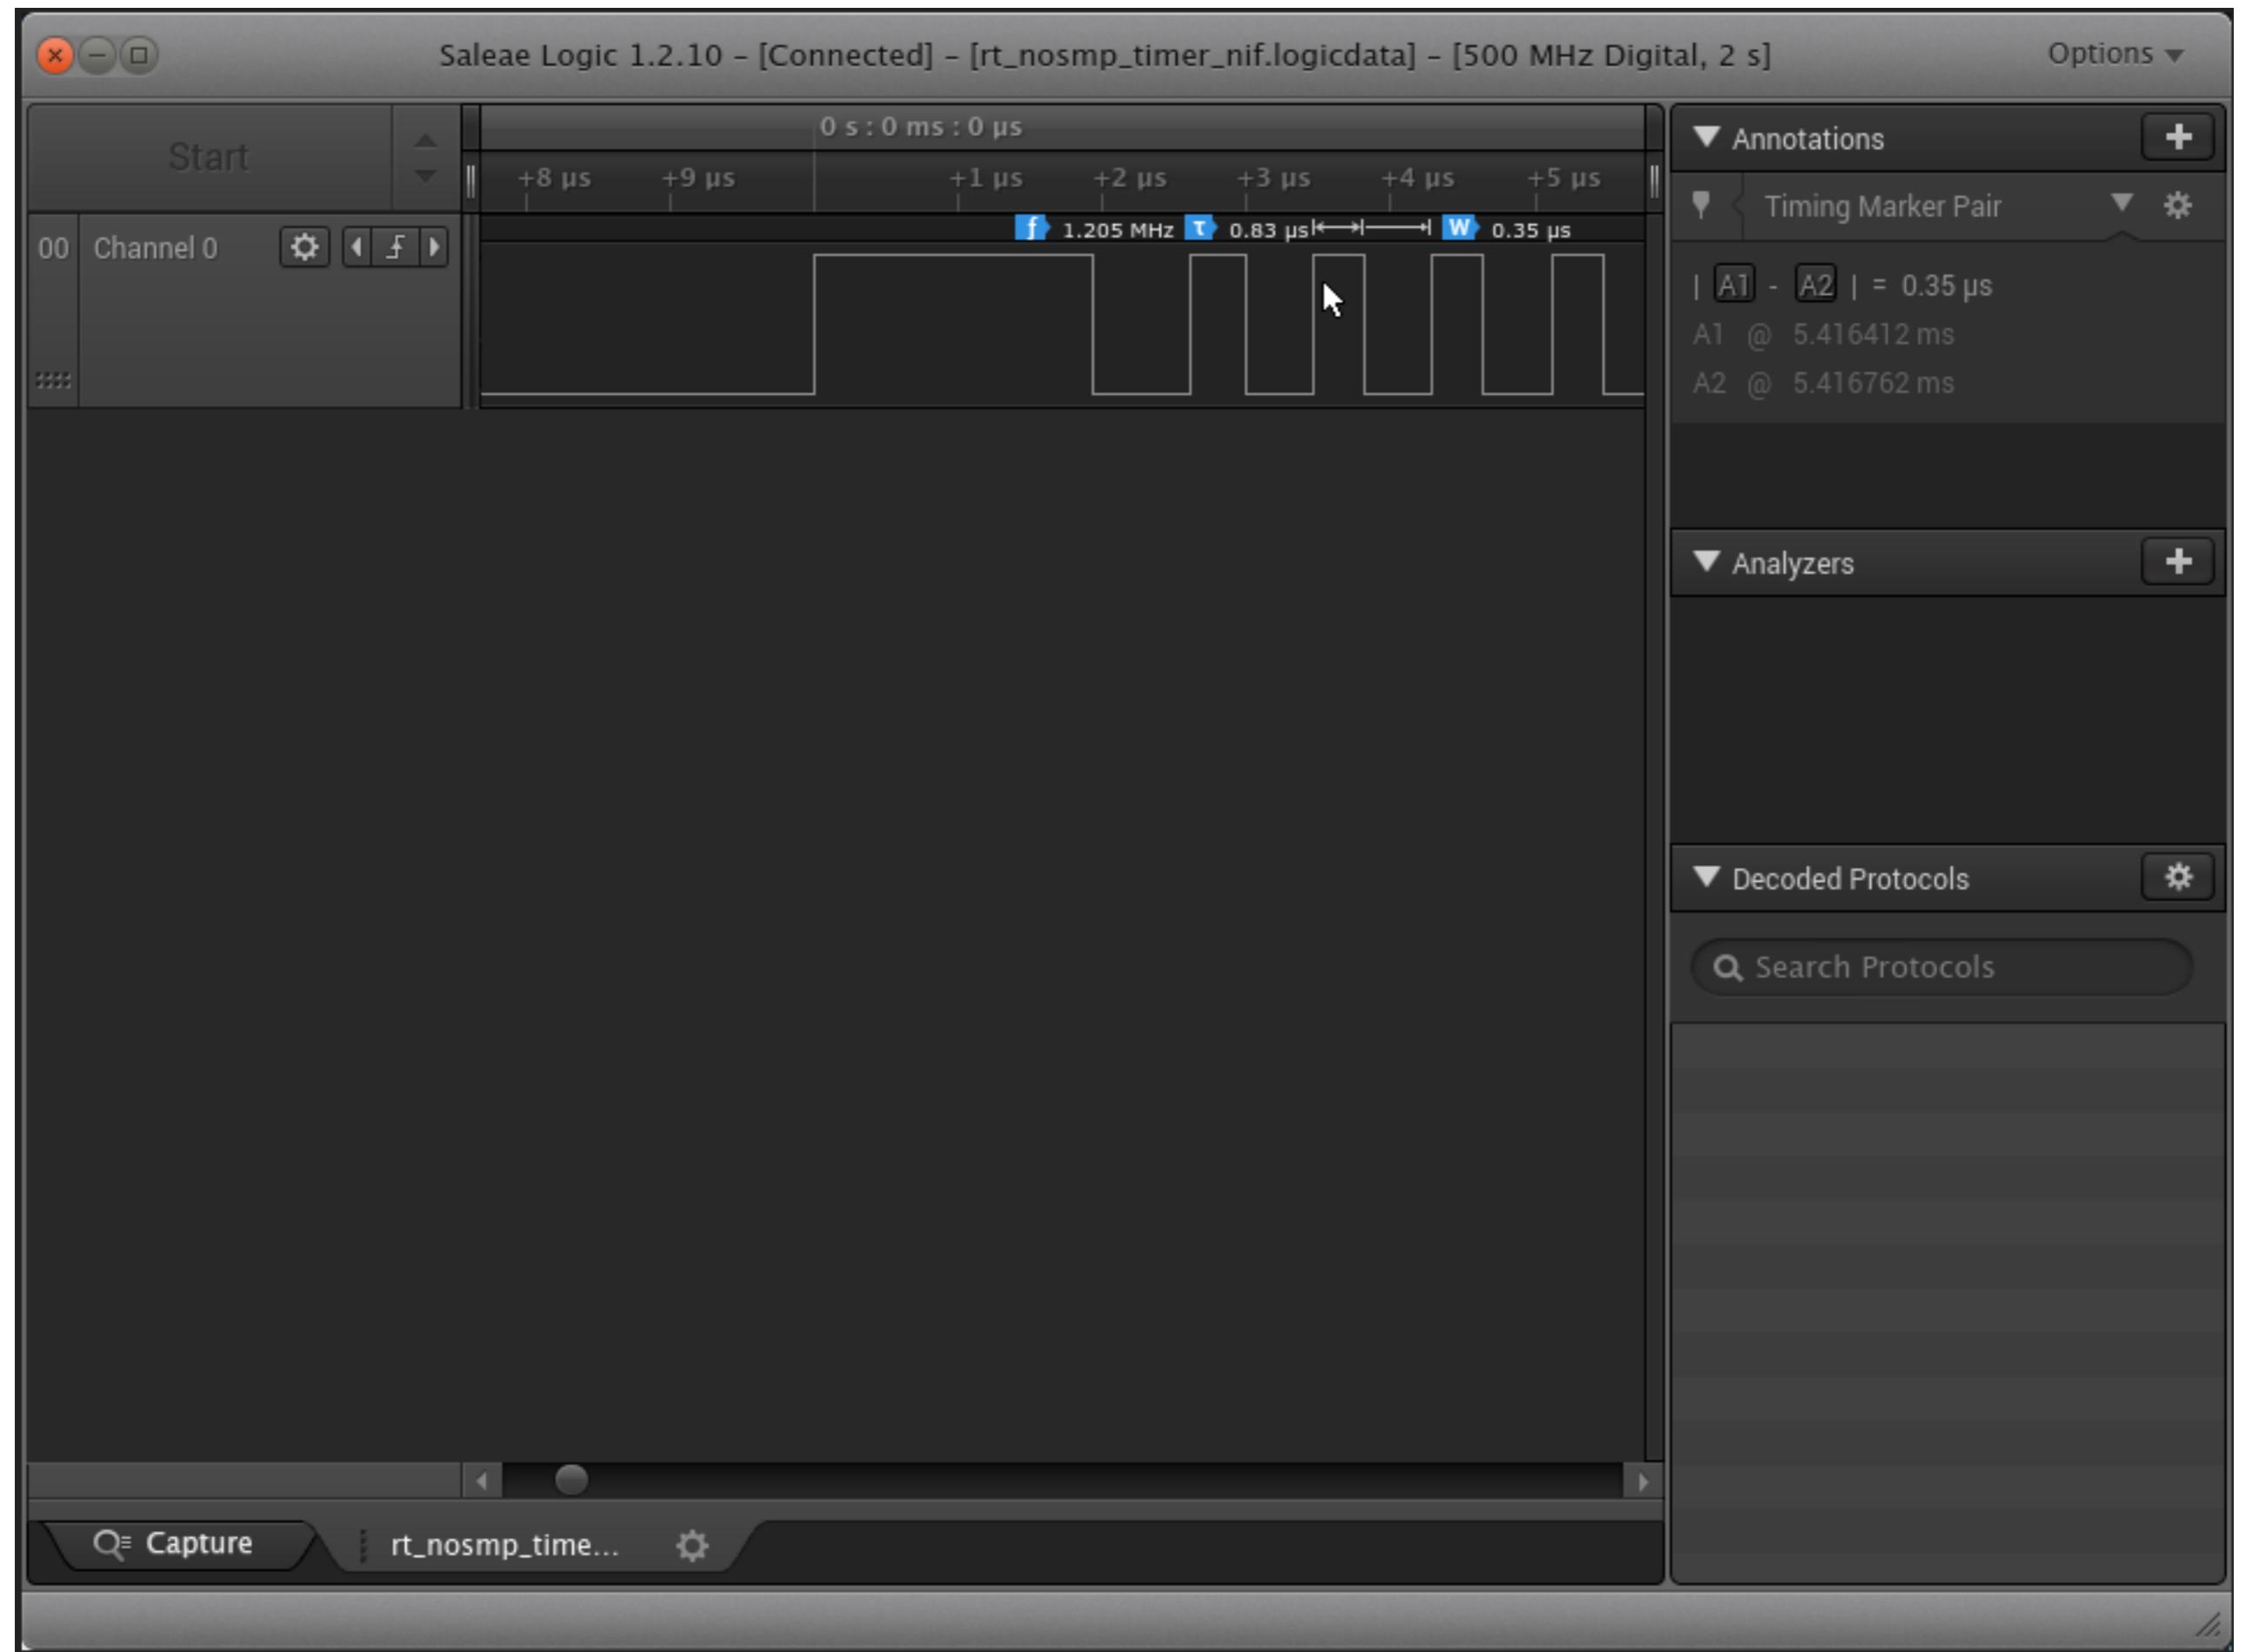
# Interfacing to C with NIFS

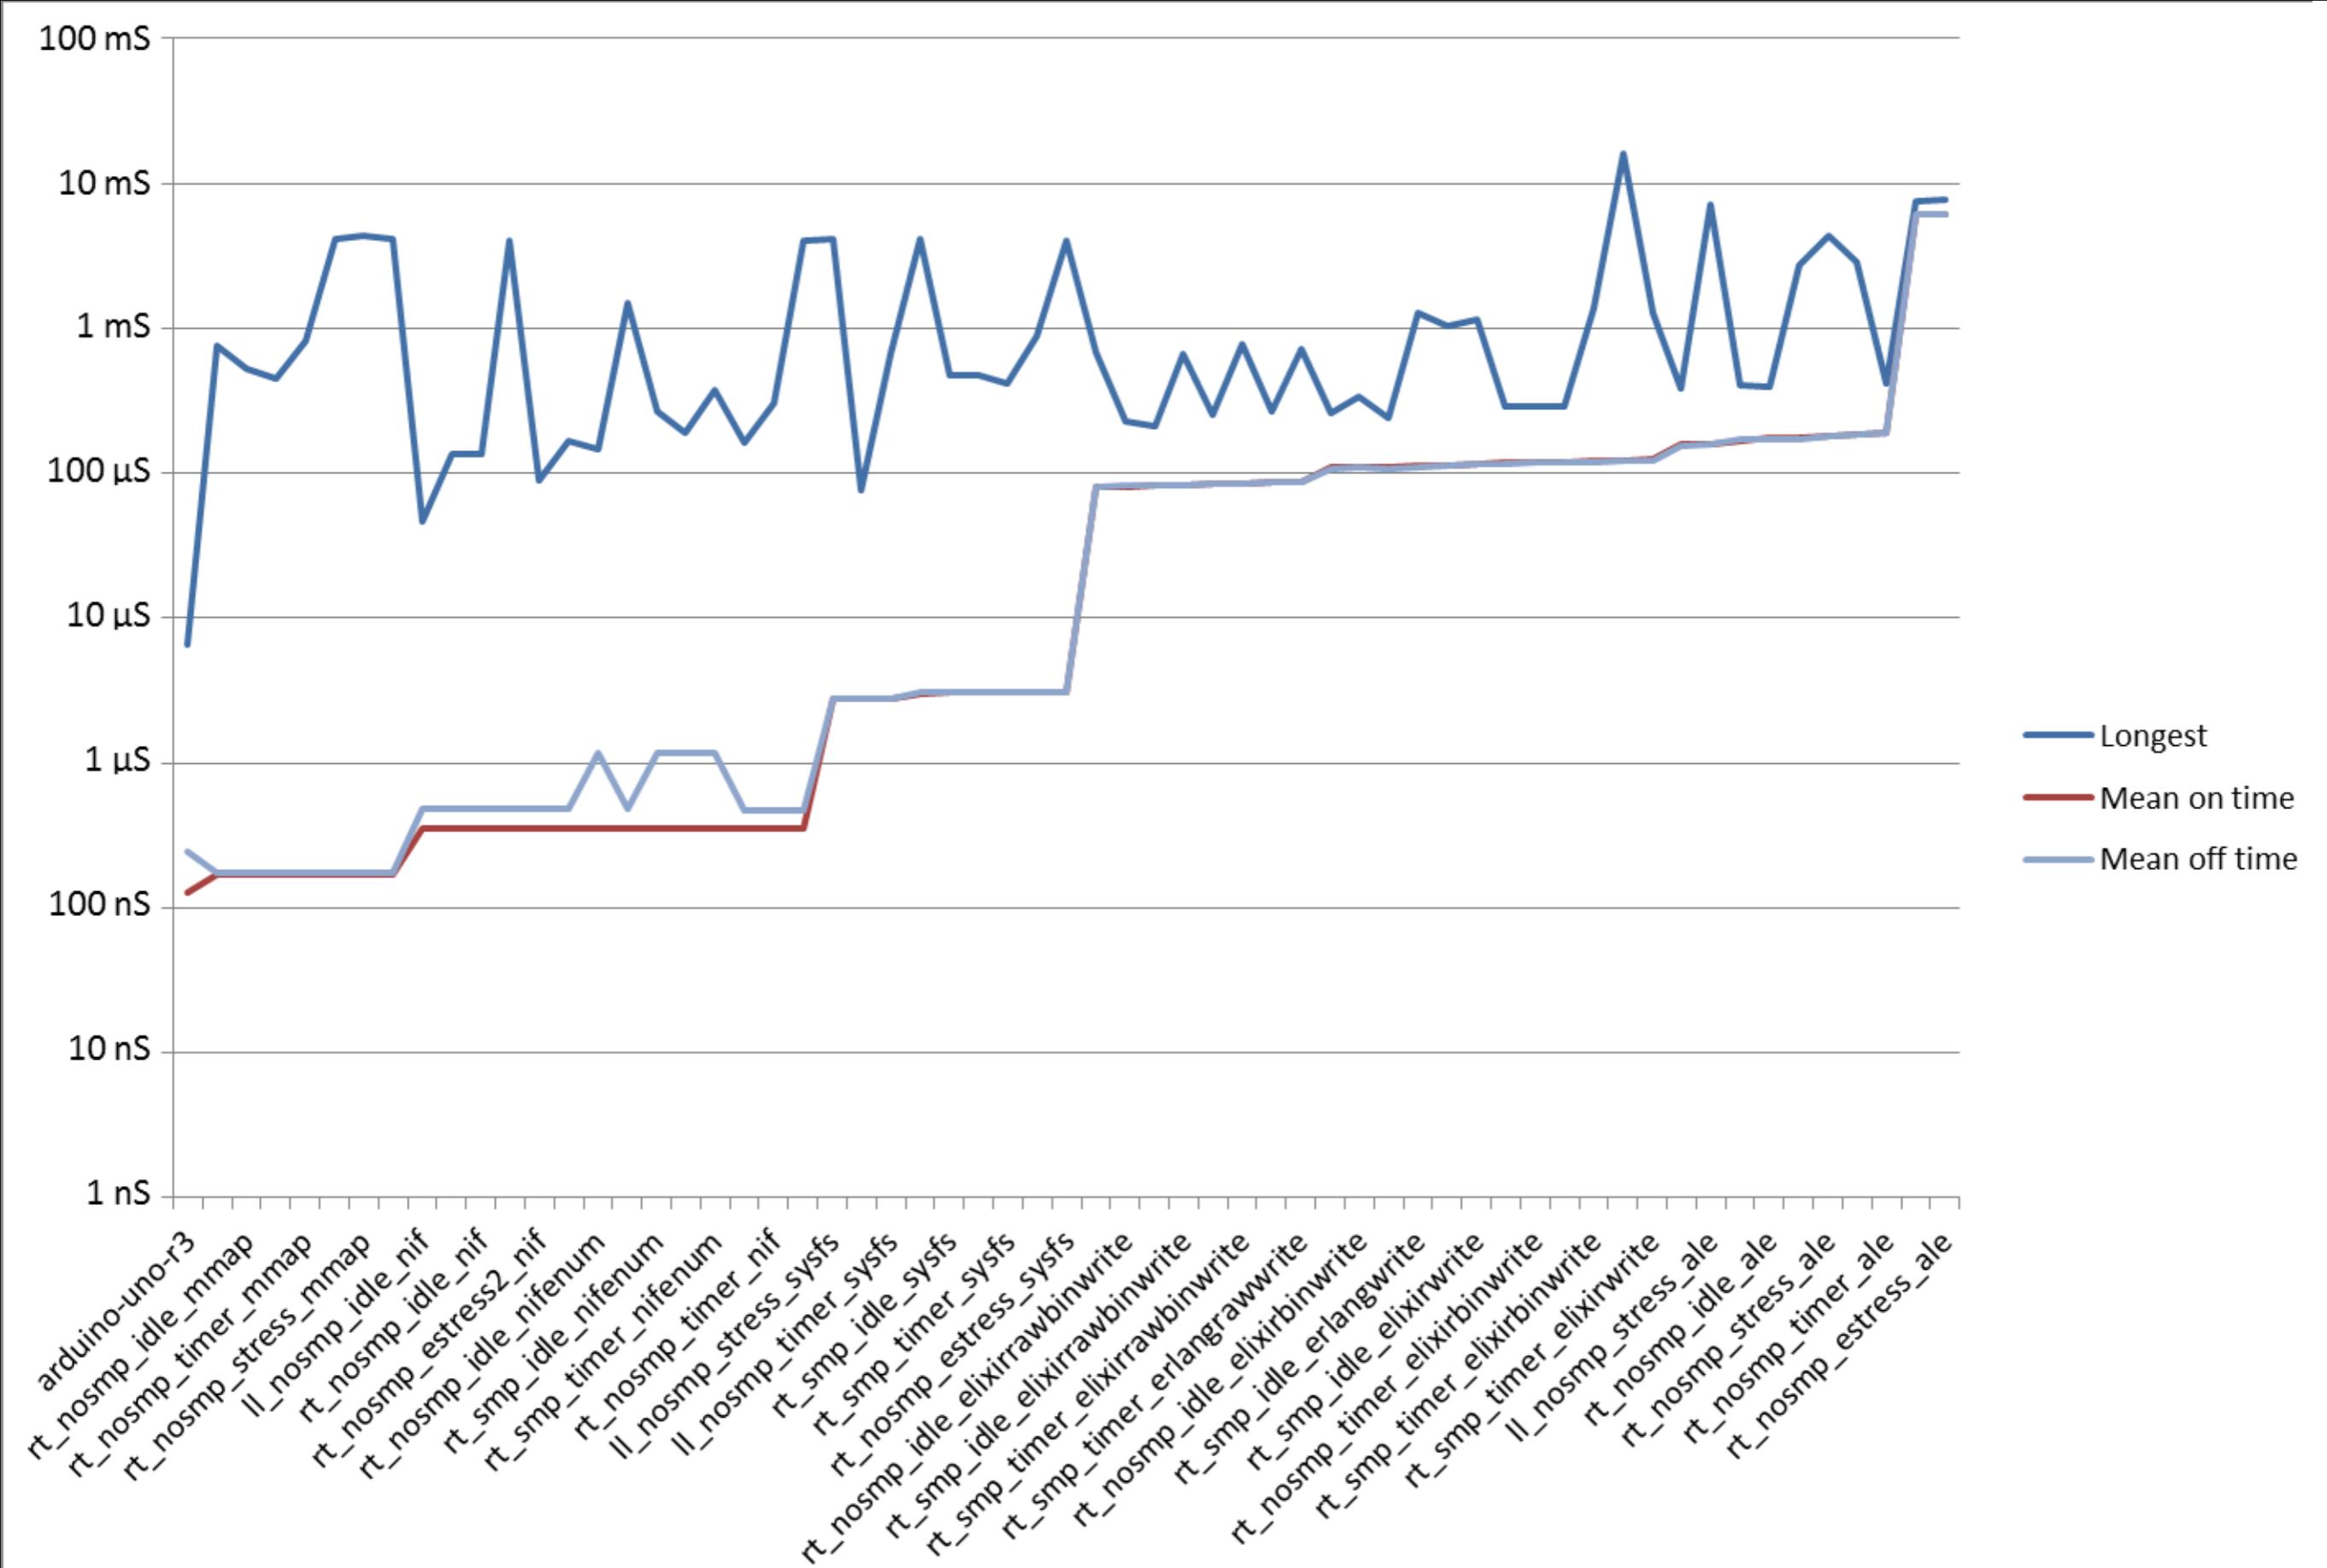
- Dynamically link a shared library into the Erlang VM
- Function callback interface in C
- Much faster than port processes
- Scheduled by the Erlang VM and run in Erlang's process environment

# Experiment

- Turn on and off an GPIO as fast as possible on a BeagleBone Black (1 GHz)
- Use a logic analyzer to measure performance
- Source code at  
[https://github.com/fhunleth/  
gpio\\_twiddler](https://github.com/fhunleth/gpio_twiddler)





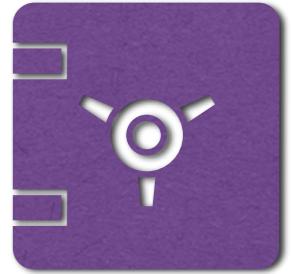


See [https://github.com/fhunleth/gpio\\_twiddle](https://github.com/fhunleth/gpio_twiddle)

# Security



# IoT Security



- Incredibly important aspect of projects today
- Developer convenience at odds with security
- Nerves is not immune to vulnerabilities

# Helpful Nerves features

- Starts small
  - Official Nerves systems have few Linux applications
  - Distillery only includes necessary OTP applications
- Read-only root filesystem\*
- Image-based firmware updates
- Erlang VM and Elixir language
- Tracks Buildroot for security patches

\*SquashFS supports appending to filesystems to add files

# Be aware

- Everything in Nerves runs as root (change in *erlinit.config*)
- IEx console available if you have physical access (change in *erlinit.config* or *vm.args*)
- Erlang distribution is not secure by default
- os:cmd/1 and ports let you get to C and Busybox easily
- No segregation of code in an Erlang VM
- Do not store secrets in firmware bundles (provision on devices)
- No support for processors that have a secure boot feature

# Security Takeaways

- Elixir and Nerves aren't magic bullets on security
- Standard recommendations apply
  - Reduce attack surface
  - Review library and application dependencies
  - Keep up-to-date with security patches
  - Secure your firmware update mechanism
  - Be mindful of transmitting and storing secrets

# Going forward

# Nerves @ ElixirConf 2017

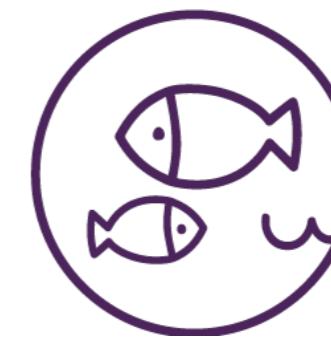
## Nerves-day (aka Thursday)

9:10 AM - 10:10 AM	Justin Schneck	Keynote
<b>Evergreen E</b> 10:40 AM - 11:20 AM	Boyd Multerer	Elixir Native UI
<b>Grand</b> 11:25 AM - 12:05 PM	Tim Mecklem	Building an Artificial Pancreas with Elixir and Nerves
<b>Evergreen E</b> 1:30 PM - 2:10 AM	Chris Côté	Embedded Elixir for Monitoring the Built Environment
<b>Grand</b> 4:10 PM - 4:50 PM	Jeff Smith	Keep an Eye on the Sky with Nerves and Phoenix

# Stay in touch!

- I'm @funkle on Slack, Twitter, Github and pretty much everywhere else
- Twitter: @NervesProject or #NervesProject
- Stop by say hi to the Nerves team in between talks
- Commercial support? Email me at [funkle@troodon-software.com](mailto:funkle@troodon-software.com)

# ELIXIRCONF™ 2017



Thanks!

# ELIXIRCONF™ 2017



## Nerves Training Part 7: Advanced Topics